


Deep Learning for System Identification

Introduction

Marco Forgione

IDSIA USI-SUPSI, Lugano, Switzerland

Artificial Intelligence, Machine Learning, Deep Learning




Machine learning: Regression

Nothing really special...

- Dataset: $D = \{(x_i, y_i)\}_{i=1}^N$
- Model structure: $\hat{y} = M(x; \theta)$
- Loss: $\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \|y_i - \hat{y}_i(\theta)\|^2$

$$\hat{\theta} = \arg \min_{\theta} \mathcal{L}(\theta)$$




Machine learning: Regression

Nothing really special...

- Dataset: $D = \{(x_i, y_i)\}_{i=1}^N$
- Model structure: $\hat{y} = M(x; \theta)$
- Loss: $\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \|y_i - \hat{y}_i(\theta)\|^2$

$$\hat{\theta} = \arg \min_{\theta} \mathcal{L}(\theta)$$




In deep learning:

- Expressive model structures like **neural networks**
- Iterative optimization, often gradient-based
- Efficient software with support for **automatic differentiation**

Is it just a form of *glorified curve fitting*?

Feedforward Neural Networks




$$z_1^2 = \sigma \left(\sum_{j=1}^4 w_{1,j} z_j^1 + b_1^1 \right)$$

$$z_2^3 = \sigma \left(\sum_{j=1}^5 w_{2,j} z_j^2 + b_2^2 \right)$$

$$y = z^4 = \sum_{j=1}^3 w_{1,j} z_j^3 + b^3$$


Feedforward Neural Networks



$$z_1^2 = \sigma \left(\sum_{j=1}^4 w_{1,j} z_j^1 + b_1^1 \right)$$
$$z_2^3 = \sigma \left(\sum_{j=1}^5 w_{2,j} z_j^2 + b_2^2 \right)$$
$$y = z^4 = \sum_{j=1}^3 w_{1,j}^3 z_j^3 + b^3$$

$$y = W_3 \sigma(W_2 \sigma(W_1 x + b_1) + b_2) + b_3 = \text{FF}(x; \theta)$$

Feedforward Neural Networks




$$z_1^2 = \sigma \left(\sum_{j=1}^4 w_{1,j} z_j^1 + b_1^1 \right)$$
$$z_2^3 = \sigma \left(\sum_{j=1}^5 w_{2,j} z_j^2 + b_2^2 \right)$$
$$y = z^4 = \sum_{j=1}^3 w_{1,j}^3 z_j^3 + b^3$$

$$y = W_3 \sigma(W_2 \sigma(W_1 x + b_1) + b_2) + b_3 = \text{FF}(x; \theta)$$

- Linear blocks interleaved by non-linear **activation functions** σ , e.g. tanh

Feedforward Neural Networks



$$z_1^2 = \sigma \left(\sum_{j=1}^4 w_{1,j} z_j^1 + b_1^1 \right)$$
$$z_2^3 = \sigma \left(\sum_{j=1}^5 w_{2,j} z_j^2 + b_2^2 \right)$$
$$y = z^4 = \sum_{j=1}^3 w_{1,j}^3 z_j^3 + b^3$$

$$y = W_3 \sigma(W_2 \sigma(W_1 x + b_1) + b_2) + b_3 = \text{FF}(x; \theta)$$

- Linear blocks interleaved by non-linear **activation functions** σ , e.g. tanh
- Non-linearities essential for **expressiveness**

Gradient Descent

1. Initialize: $\hat{\theta}$
2. **for** $k = 1, \dots, N$:
 - Compute Gradient: $g = \nabla_{\theta} \mathcal{L}(\theta^k)$
 - Update Parameters:

$$\hat{\theta} = \hat{\theta} - \gamma g$$

3. **end**


Gradient Descent

1. Initialize: $\hat{\theta}$
2. **for** $k = 1, \dots, N$:
 - Compute Gradient: $g = \nabla_{\theta} \mathcal{L}(\theta^k)$
 - Update Parameters:

$$\hat{\theta} = \hat{\theta} - \gamma g$$

3. **end**

Local convergence to one of the **several minima** is OK!



Gradient Descent


1. Initialize: $\hat{\theta}$
2. **for** $k = 1, \dots, N$:
 - Compute Gradient: $g = \nabla_{\theta} \mathcal{L}(\theta^k)$
 - Update Parameters:

$$\hat{\theta} = \hat{\theta} - \gamma g$$

3. **end**

Local convergence to one of the **several minima** is OK!

Loss landscape




- Mini-batching: at each iteration, compute loss on a subset of the dataset
- Variants of plain gradient descent like Adam are more common and effective
- Second-order methods like (L)-BFGS also suitable for small/medium-scale problems

Example: synthetic toy dataset

Consider the 2D function $f : \mathbb{R}^2 \mapsto \mathbb{R}$

$$f(x) = 2 \sin(x_1) - 3 \cos(x_2)$$
$$x \in [-2, 2]^2 \subset \mathbb{R}^2$$

- Training and test datasets: 500 random points in the domain
- Additive noise with standard deviation 0.1





Feed-forward neural network

Model with:

- 2 inputs, 1 output - just like $f(x)$
- 2 hidden layers with [16, 8] neurons
- $\tanh(\cdot)$ activation functions

$$\hat{y} = W_3 \tanh(W_2 \tanh(W_1 x + b_1) + b_2) + b_3$$

Training: 5000 iterations of Adam...




Model evaluation

Common to inspect on the test:

- Predictions and residuals vs measured output (left)
- Histogram of residuals (right)

Keep on following the good practices!



Model evaluation


Common to inspect on the test:

- Predictions and residuals vs measured output (left)
- Histogram of residuals (right)

Keep on following the good practices!

2D toy example. We can also visualize:

- The known true function (right)
- The learned model over a grid (left)



Deep Learning For System Identification

- Feedforward nets fed by lagged input/outputs for NARX identification:

$$\hat{y}(k) = \text{FF}(\overbrace{y(k-1), \dots, y(k-n_a), u(k), u(k-1), \dots, u(k-n_b-1)}^{=x(k)}; \theta)$$

Deep Learning For System Identification

- Feedforward nets fed by lagged input/outputs for NARX identification:

$$\hat{y}(k) = \text{FF}\left(\overbrace{y(k-1), \dots, y(k-n_a), u(k), u(k-1), \dots, u(k-n_b-1)}^{=x(k)}; \theta\right)$$

- State-space models with neural-network update/output maps:
 - Also known as Recurrent Neural Networks

$$\begin{aligned}x(k+1) &= \text{FF}_x(x(k), u(k); \theta) \\y(k) &= \text{FF}_y(x(k); \theta)\end{aligned}$$

What we won't cover

- Advanced architectures
 - physics-informed
 - grounded on system theory (e.g., stable by design)
 - transformers, diffusion, foundation models,...

What we won't cover


- Advanced architectures
 - physics-informed
 - grounded on system theory (e.g., stable by design)
 - transformers, diffusion, foundation models,...
- Advanced optimization
 - mini-batching (to handle large datasets)
 - second-order, non-smooth, constrained...
 - automatic differentiation details

What we won't cover

- Advanced architectures
 - physics-informed
 - grounded on system theory (e.g., stable by design)
 - transformers, diffusion, foundation models,...
- Advanced optimization
 - mini-batching (to handle large datasets)
 - second-order, non-smooth, constrained...
 - automatic differentiation details
- Theoretical aspects
 - learning theory
 - uncertainty quantification
 - ...

Software for Deep Learning

An essential aspect is the software implementation. A few options:



- PyTorch and JAX: libraries on top of Python
- Julia: a programming language focused on numerical computations
- MATLAB: you should know it already...

Software for Deep Learning

An essential aspect is the software implementation. A few options:



- PyTorch and JAX: libraries on top of Python
- Julia: a programming language focused on numerical computations
- MATLAB: you should know it already...

We will try out Google's **JAX**!

Some Literature

Machine learning reference books

- Murphy, Kevin P. *Probabilistic machine learning: An introduction*. MIT press, 2022.
- Murphy, Kevin P. *Probabilistic machine learning: Advanced topics*. MIT press, 2023

Deep learning in system identification (my biased perspective)

- Forggione, Marco, and Piga, Dario. "Continuous-time system identification with neural networks: Model structures and fitting criteria." *European Journal of Control* 59 (2021).
- Beintema, Gerben I., Maarten Schoukens, and Roland Tóth. "Deep subspace encoders for nonlinear system identification." *Automatica* (2023).
- Bemporad, Alberto. "An L-BFGS-B Approach for Linear and Nonlinear System Identification Under ℓ_1 and Group-Lasso Regularization." *IEEE Transactions on Automatic Control* (2025).
- Pillonetto, Gianluigi, Aleksandr Aravkin, Daniel Gedon, Lennart Ljung, Antonio H. Ribeiro, and Thomas B. Schön. "Deep networks for system identification: a survey." *Automatica* (2025).