

# Deep Learning for System Identification

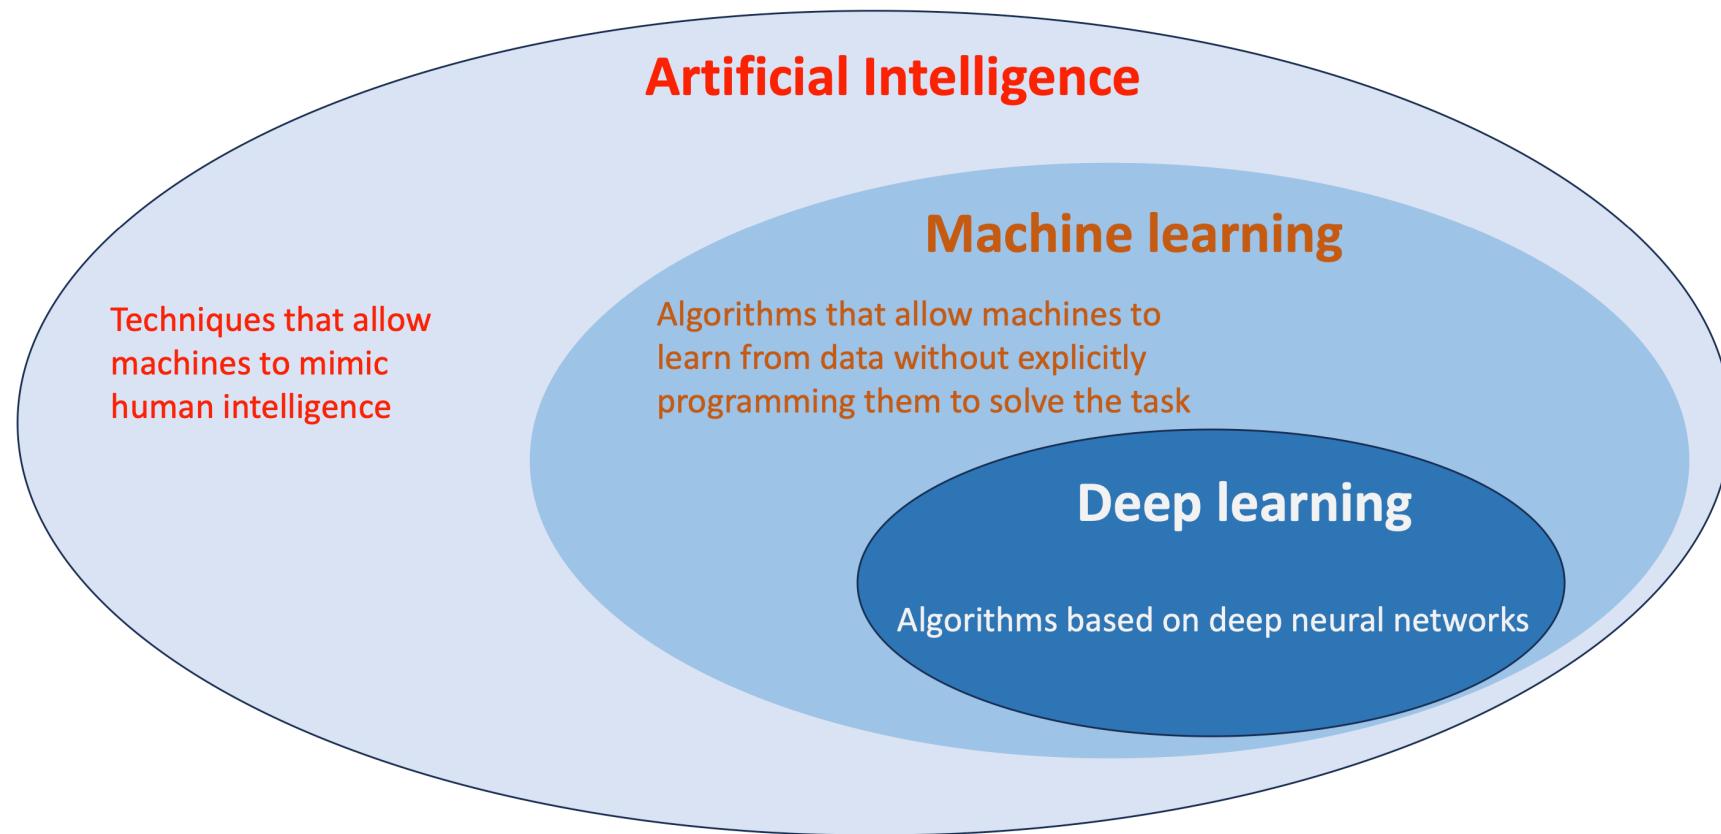
Marco Forgione

SUPSI – Scuola Universitaria Professionale della Svizzera Italiana

IDSIA – Dalle Molle Institute for Artificial Intelligence

[marco.forgione@supsi.ch](mailto:marco.forgione@supsi.ch)

# Artificial Intelligence, Machine Learning, Deep Learning



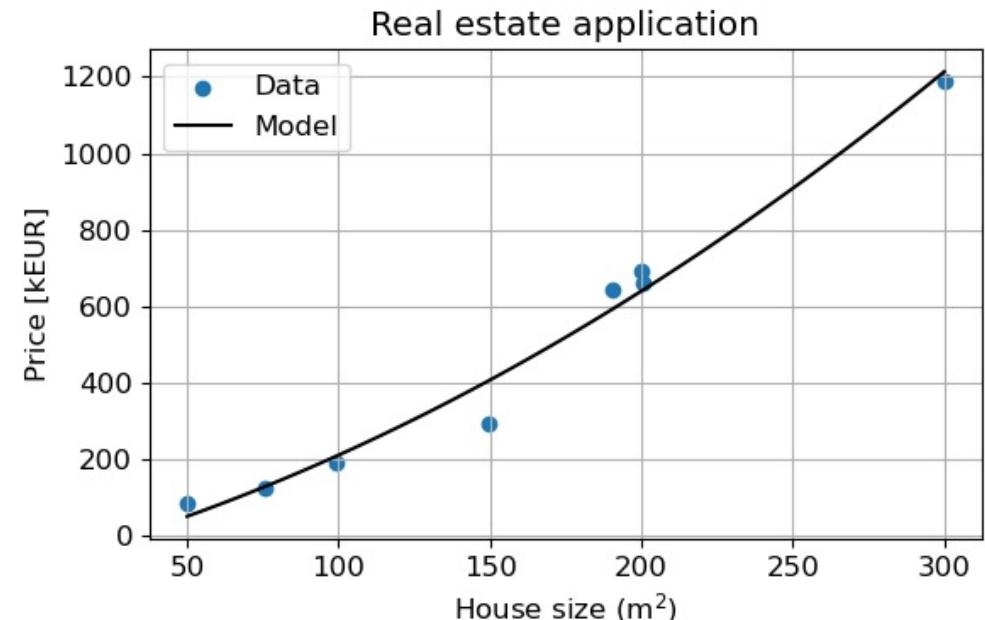
## Machine learning: Regression

- Dataset:  $D = \{(x_i, y_i)\}_{i=1}^N$

- Model structure:  $\hat{y} = M(x; \theta)$

- Loss:  $\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N \|y_i - \hat{y}_i(\theta)\|^2$

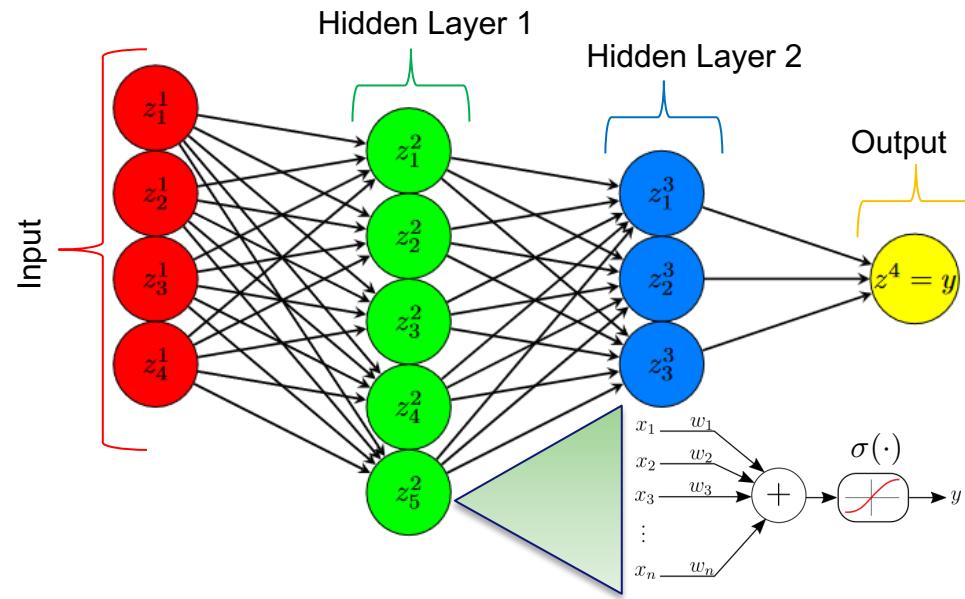
$$\hat{\theta} = \arg \min_{\theta} \mathcal{L}(\theta)$$



- Expressive model structures like neural networks
- Iterative optimization, often gradient-based
- Efficient software for automatic differentiation
- Still, it is just *glorified curve fitting*...

# Feed-Forward Neural Networks

Just a particular model structure



$$z_1^2 = \sigma \left( \sum_{j=1}^4 w_{1,j} z_j^1 + b_1 \right)$$

$$z_2^3 = \sigma \left( \sum_{j=1}^5 w_{2,j} z_j^2 + b_2 \right)$$

$$y = z^4 = \sum_{j=1}^3 w_{3,j} z_j^3 + b^3$$

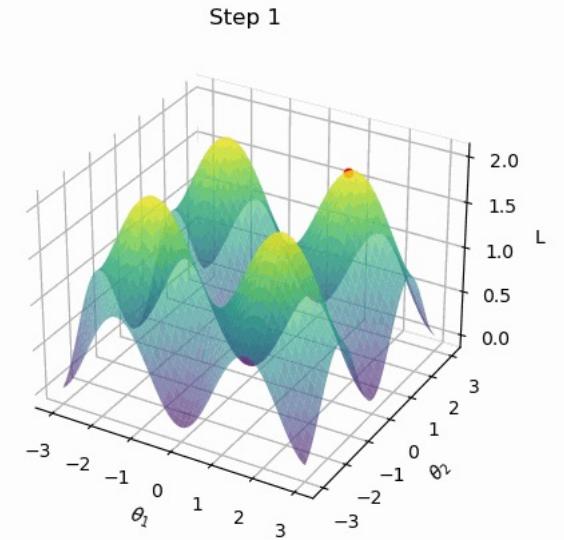
$$y = W_3 \sigma(W_2 \sigma(W_1 x + b_1) + b_2) + b_3 = \text{FF}(x; \theta)$$

- Linear blocks interleaved by element-wise **non-linearities** (activation functions)
- Non-linearities essential for expressiveness

# Gradient Descent

1. Initialize Parameters:  $\theta^{(0)}$
2. for  $k = 0, 1, \dots$ : until maximum number of iterations or convergence do:
  - (a) Compute Gradient:  $\nabla_{\theta} \mathcal{L}(\theta^{(k)})$
  - (b) Update Parameters:

$$\theta^{(k+1)} = \theta^{(k)} - \gamma \cdot \nabla_{\theta} \mathcal{L}(\theta^{(k)})$$



Local convergence to one of the **several minima** is OK!

- Mini-batching: at each iteration, compute loss & gradients on a subset of the dataset
- Variants of plain gradient descent like Adam are now more common
- Second-order methods like (L)-BFGS suitable for problems up to medium scale
  - most sysid and scientific machine-learning problems

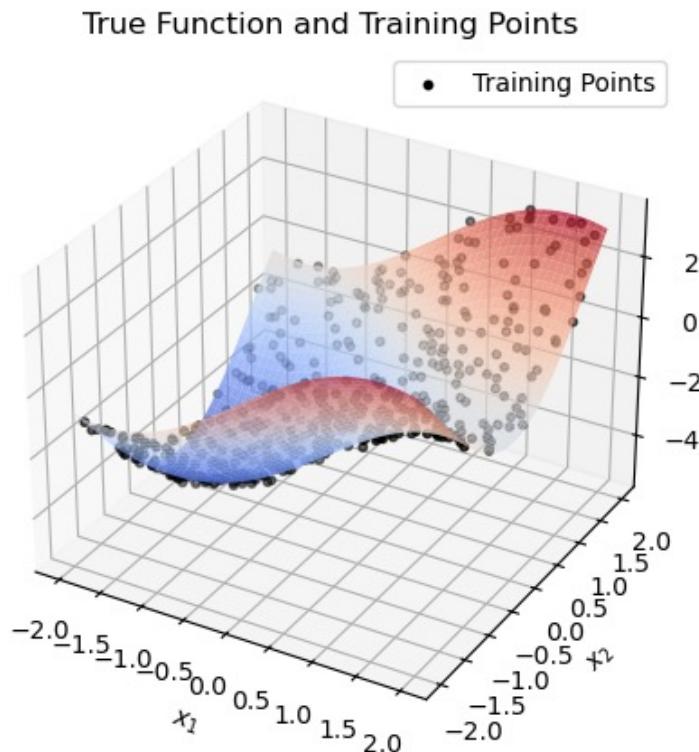
## Example: synthetic toy dataset

Consider the 2D function  $f : \mathbb{R}^2 \rightarrow \mathbb{R}$

$$f(x) = 2 \sin(x_1) - 3 \cos(x_2)$$

$$x \in [-2, 2]^2 \subset \mathbb{R}^2$$

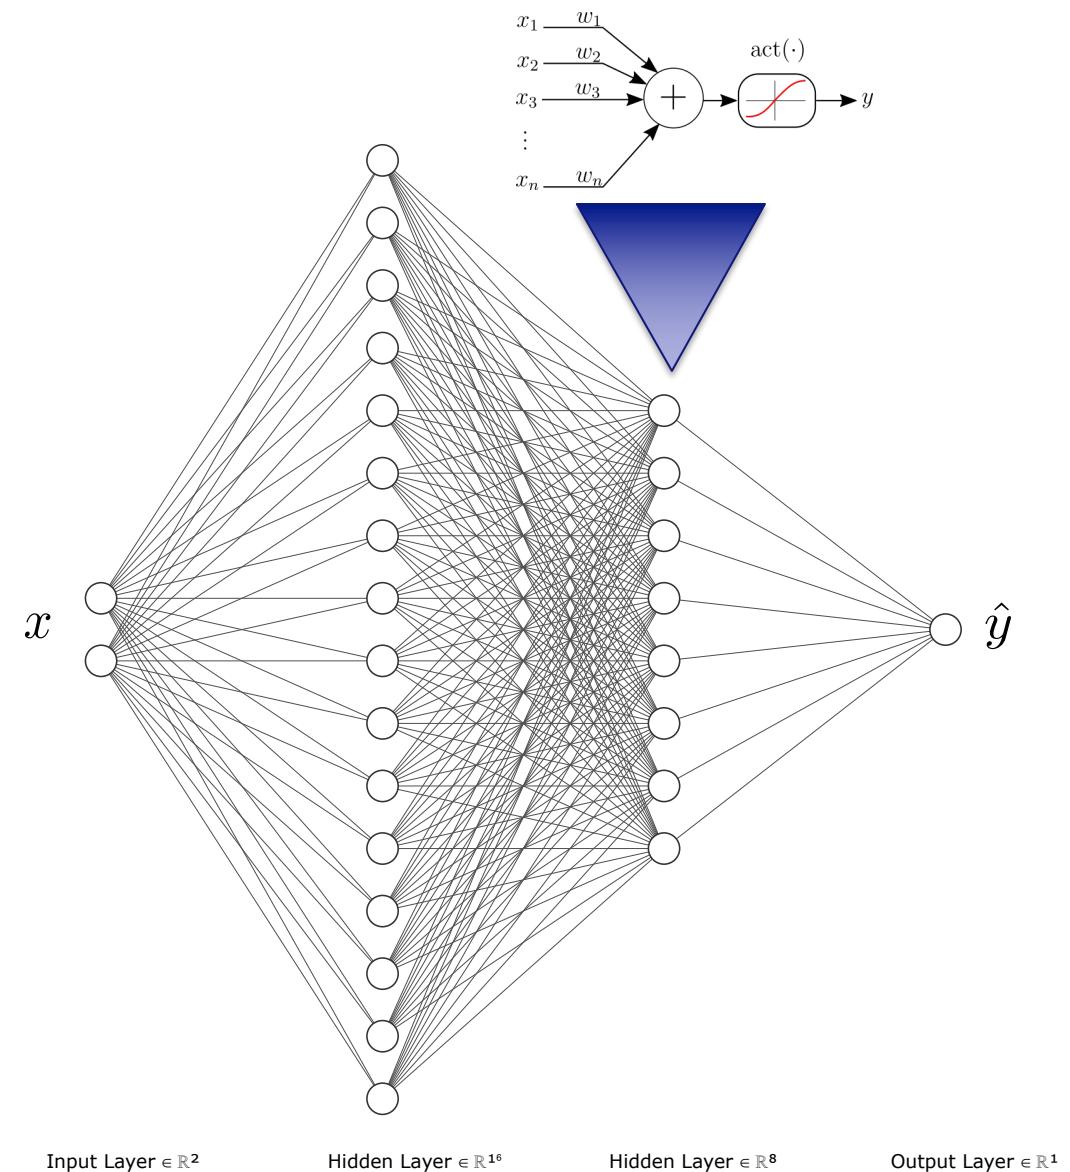
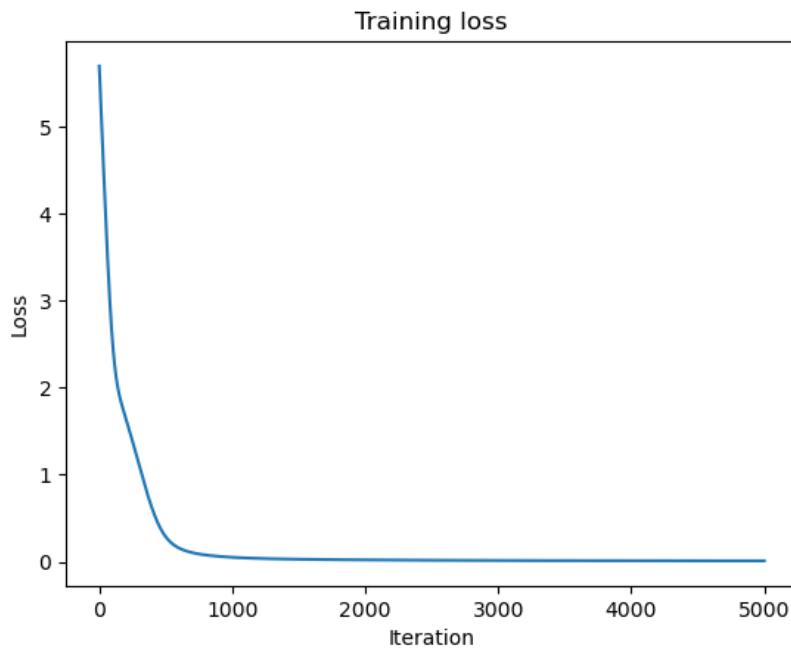
- Training and test datasets: 500 points uniformly sampled in the domain.
- Additive noise with standard deviation 0.1



## Feed-forward neural network

- 2 inputs, 1 output - this is the structure of  $f(x)$
- 2 hidden layers with [16, 8] neurons
- tanh non-linearities

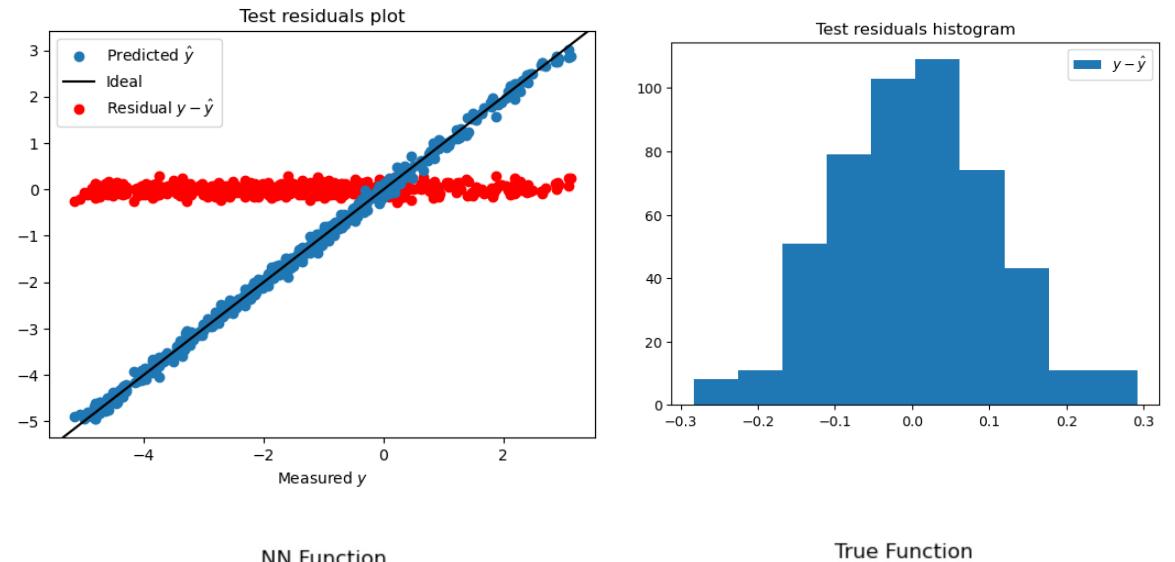
Training with 5000 iterations of Adam...



# Model evaluation

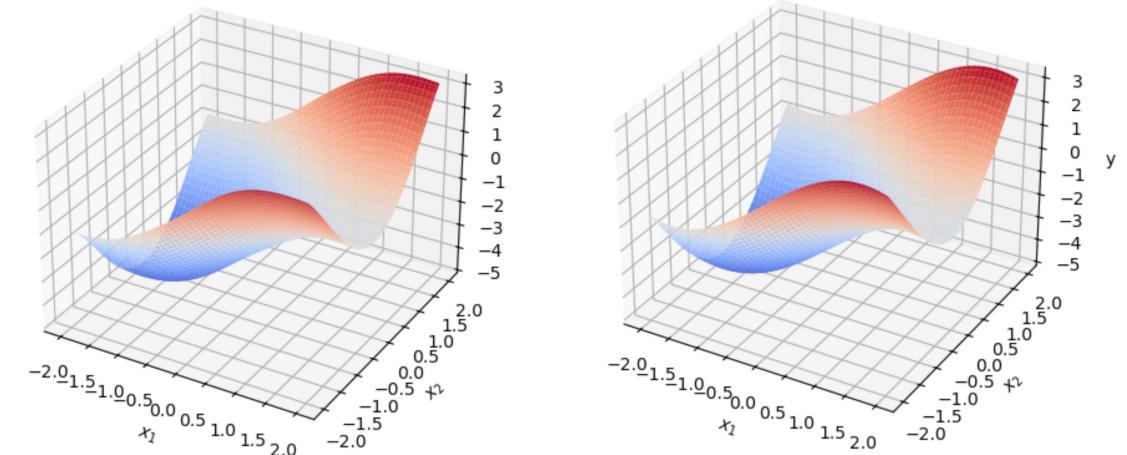
It is common to inspect on the test dataset:

- Predictions and residuals vs measured output (left)
- Histogram of residuals (right)
- Keep on following good statistical practices!



In this 2D toy example, we can also visualize:

- The learned function over a grid (left)
- The known true function (right)



## Deep Learning for System Identification

- Feed-forward nets fed by lagged input/outputs are directly applicable for NARX regression

$$\hat{y}(k) = \text{FF}\left(\overbrace{y(k-1), \dots, y(k-n_a), u(k), u(k-1), \dots, u(k-n_b-1)}^{=x(k)}; \theta\right)$$

- State-space models with neural-network update/output maps:
  - Also known as Recurrent Neural Networks

$$x(k+1) = \text{FF}_x(x(k), u(k); \theta)$$

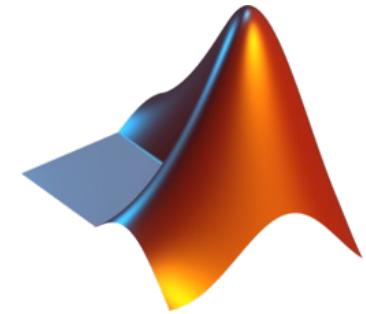
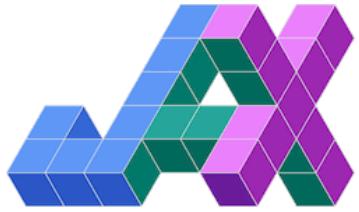
$$y(k) = \text{FF}_y(x(k); \theta)$$

## What we won't cover

- Advanced architectures
  - physics-informed
  - grounded on system theory (e.g., stable by design)
  - transformers, diffusion, foundation models,...
- Optimization
  - mini-batching (to handle large datasets)
  - second-order, non-smooth, constrained...
  - automatic differentiation details
- Theoretical aspects
  - learning theory
  - uncertainty quantification
  - ...

# Software for Deep Learning

An essential aspect is the software implementation. Non-exhaustive options are:



- PyTorch and JAX: libraries on top of Python
- Julia: a programming language focused on numerical computations
- MATLAB: you should know it already...

We will try out **JAX**!

## Literature

### A machine learning reference book

- Murphy, Kevin P. *Probabilistic machine learning: an introduction*. MIT press, 2022.
- Murphy, Kevin P. *Probabilistic machine learning: Advanced topics*. MIT press, 2023

### Deep learning in system identification (my biased perspective)

- Forgione, Marco, and Dario Piga. "Continuous-time system identification with neural networks: Model structures and fitting criteria." *European Journal of Control* 59 (2021): 69-81.
- Beintema, Gerben I., Maarten Schoukens, and Roland Tóth. "Deep subspace encoders for nonlinear system identification." *Automatica* 156 (2023): 111210.
- Bemporad, Alberto. "An L-BFGS-B Approach for Linear and Nonlinear System Identification Under  $\ell_1$  and Group-Lasso Regularization." *IEEE Transactions on Automatic Control* (2025).
- Pillonetto, Gianluigi, Aleksandr Aravkin, Daniel Gedon, Lennart Ljung, Antonio H. Ribeiro, and Thomas B. Schön. "Deep networks for system identification: a survey." *Automatica* 171 (2025): 111907.